

D213_Time_Series

April 4, 2022

```
[1]: # Standard data science imports
import numpy as np
import pandas as pd
# Visualization libraries
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: # Set plot to ggplot for appearance
plt.style.use('ggplot')
```


```
[3]: # Load data set into Pandas dataframe
df_teleco=pd.read_csv('c:/Users/almingah/Desktop/MSDA-WGU/D213/
↳teleco_time_series.csv')
```

```
[4]: df_teleco.head()
```

```
[4]:
```

	Day	Revenue
0	1	0.000000
1	2	0.000793
2	3	0.825542
3	4	0.320332
4	5	1.082554

```
[5]: # line graph visualizing the realization of the time series
plt.figure(figsize=(18, 6))
plt.plot( df_teleco['Day'], df_teleco['Revenue'], color='tab:blue')
plt.xlabel('Time(Day)')
plt.ylabel('Revenue in millions, USD')
plt.title(' Revenue for The last two years')
plt.show()
```



```
[6]: # Get dataset size
df_teleco.shape
```

```
[6]: (731, 2)
```

```
[7]: # view statistics of dataset
df_teleco.describe()
```

```
[7]:
```

	Day	Revenue
count	731.000000	731.000000
mean	366.000000	9.822901
std	211.165812	3.852645
min	1.000000	0.000000
25%	183.500000	6.872836
50%	366.000000	10.785571
75%	548.500000	12.566911
max	731.000000	18.154769

```
[8]: # Get info
df_teleco.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Day      731 non-null      int64
1   Revenue  731 non-null      float64
dtypes: float64(1), int64(1)
memory usage: 11.5 KB
```

```
[9]: # check dataset for missing data points
df_nulls = df_teleco.isnull().sum()
```

```
print(df_nulls)
```

```
Day          0  
Revenue      0  
dtype: int64
```


```
[10]: # check for duplicates in Day  
print(df_teleco.Day.duplicated().sum())
```

```
0
```

```
[11]: #Check for NAs  
print(df_teleco.isna().sum())
```

```
Day          0  
Revenue      0  
dtype: int64
```

```
[12]: # Visualize missing values in dataset (GeeksForGeeks, p. 1)  
import missingno as msno  
msno.matrix(df_teleco);
```



```
[13]: # Import adfuller and order libraries  
from statsmodels.tsa.stattools import adfuller  
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf  
import statsmodels.tsa.stattools as ts  
from statsmodels.tsa.seasonal import seasonal_decompose  
import statsmodels.api as sm  
from sklearn.metrics import mean_squared_error  
from statsmodels.tools.eval_measures import rmse
```


```
[14]: # instantiate adfuller (VERMA, Y. 2021)
adft = adfuller(df_teleco.iloc[:, 1].values, autolag='AIC')
print("1. ADF c-value : ",adft[0])
print("2. P-Value : ", adft[1])
print("3. Num Of Lags : ", adft[2])
print("4. Num Of Observations Used :", adft[3])
print("5. Critical Values :")
for key, val in adft[4].items():
    print("\t",key, ": ", val)
```

```
1. ADF c-value : -1.9246121573101826
2. P-Value : 0.3205728150793969
3. Num Of Lags : 1
4. Num Of Observations Used : 729
5. Critical Values :
    1% : -3.4393520240470554
    5% : -2.8655128165959236
    10% : -2.5688855736949163
```

```
[15]: #set up Stationairty test
def test_stationarity(timeseries):
    #Determing rolling statistics
    rolmean = timeseries.rolling(window=30).mean()
    rolstd = timeseries.rolling(window=30).std()

    #Plot rolling statistics:
    orig = plt.plot(timeseries, color='blue',label='Original')
    mean = plt.plot(rolmean, color='red', label='Rolling Mean')
    std = plt.plot(rolstd, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean & Standard Deviation')
    plt.show(block=False)
```

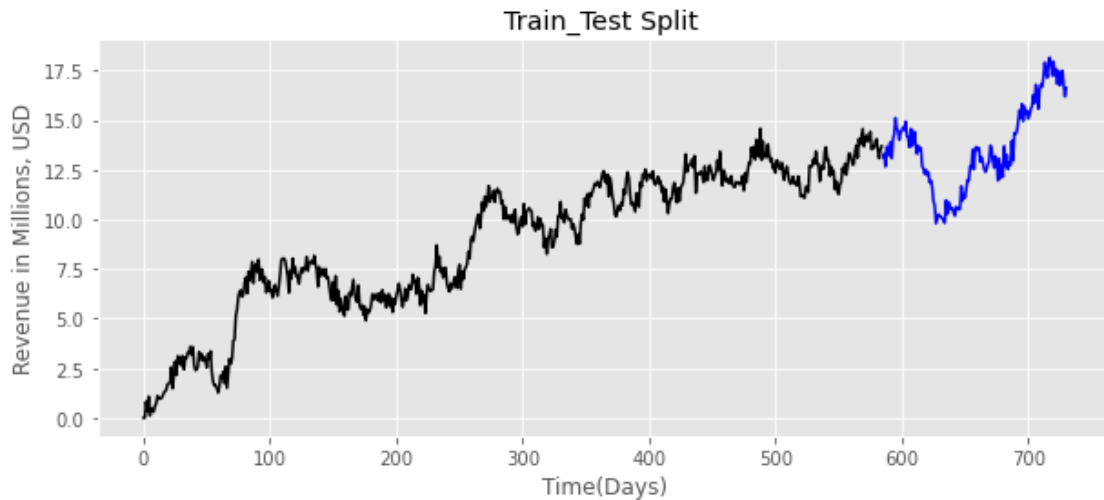
```
[17]: # perform test on entire dataset( again, the mean is not zero, hence dataset_
↳not stationary)
test_stationarity(df_teleco['Revenue'])
```



```
[18]: # split dataset into training and test sets with 80% for training( 80% of 731
      ↪ days= 585)
df_teleco['Day'] = df_teleco.index
df_train = df_teleco[:585]
```


```
[19]: # 20% Test ( from row 585 to the end of the data)
df_test = df_teleco[585:]
```

```
[20]: # Plot training and test datasets
plt.figure(figsize=(10, 4))
plt.plot(df_train['Revenue'], color='black')
plt.plot(df_test['Revenue'], color='blue')
plt.title('Train_Test Split ')
plt.xlabel('Time(Days)')
plt.ylabel('Revenue in Millions, USD')
plt.show()
```




```
[21]: # save a Clean dataset as csv
df_teleco.to_csv('teleco_time_series_clean.csv')
```


```
[22]: #decompose the training dataset
decomposition=seasonal_decompose(df_train['Revenue'], model='additive',
    ↪period=30)
trend=decomposition.trend
seasonal=decomposition.seasonal
residual=decomposition.resid
plt.subplot(411)
plt.plot(df_train['Revenue'],color='red', label='original_series')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend,color='blue', label='Trend')
plt.legend(loc='best')
plt.tight_layout()
plt.subplot(414)
plt.plot(residual,color='green', label='Residual')
plt.legend(loc='best')
plt.tight_layout()
plt.subplot(413)
plt.plot(seasonal,color='brown', label='Seasonality')
plt.legend(loc='best')
plt.tight_layout()
plt.show()
```



```
[23]: # Plot the trend component only
plt.figure(figsize=(12, 4))
plt.plot(trend, color='tab:blue')
plt.xlabel('Day')
plt.ylabel('Revenue Trend')
plt.title('Decomposed Revenue Trend')
plt.show()
```





```
[24]: #Plot autocorrelation function on training set
plot_acf(df_train['Revenue']);
```



```
[25]: #Plot partial autocorrelation function on training set
plot_pacf(df_train['Revenue'], method='ywmm')
```

[25]:



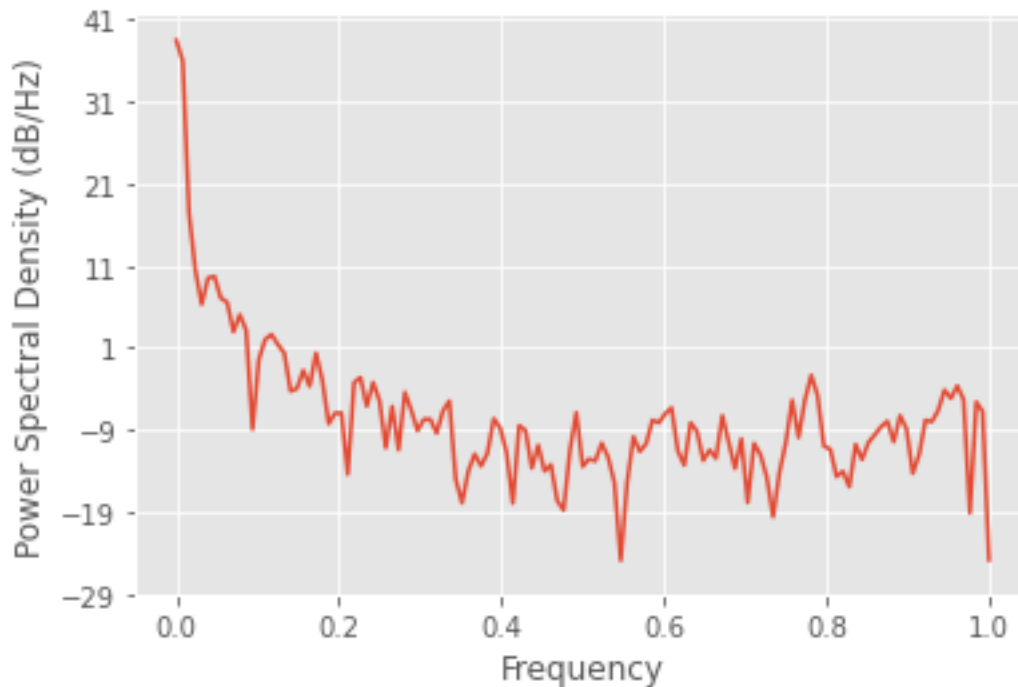
```
[26]: # Run spectral density function
plt.psd(df_teleco['Revenue'])
```

```
[26]: (array([6.97387711e+03, 3.91439441e+03, 5.47611144e+01, 1.10791953e+01,
 4.25439684e+00, 8.90767103e+00, 9.32436541e+00, 5.07699490e+00,
 4.50820186e+00, 1.97142105e+00, 3.17930644e+00, 2.05465010e+00,
 1.27856111e-01, 9.12357670e-01, 1.58928494e+00, 1.82600308e+00,
 1.37345831e+00, 1.08602805e+00, 3.71664573e-01, 4.01447712e-01,
 6.67019362e-01, 4.28134847e-01, 1.08303874e+00, 5.05461139e-01,
 1.49612021e-01, 2.01156666e-01, 2.01929017e-01, 3.62482513e-02,
 4.74386760e-01, 5.46619399e-01, 2.44460915e-01, 4.75225583e-01,
 2.85007556e-01, 7.59631735e-02, 2.43183009e-01, 7.19503505e-02,
 3.60475608e-01, 2.18971480e-01, 1.23421972e-01, 1.69234109e-01,
 1.69371921e-01, 1.14748035e-01, 2.14144408e-01, 2.82736486e-01,
 3.17059554e-02, 1.62458482e-02, 3.98324693e-02, 6.38895068e-02,
 4.57061837e-02, 6.54718394e-02, 1.72598760e-01, 1.32802101e-01,
 6.74716970e-02, 1.61089989e-02, 1.42266347e-01, 1.23640967e-01,
 4.30737122e-02, 8.17619769e-02, 3.95102804e-02, 4.74833201e-02,
 1.73085344e-02, 1.32485092e-02, 7.14595244e-02, 2.04348554e-01,
 4.51051907e-02, 5.55609932e-02, 5.26341931e-02, 8.68327071e-02,
 5.86340275e-02, 2.88124318e-02, 3.23398046e-03, 2.85277556e-02,
 1.03859619e-01, 6.85761289e-02, 8.53467489e-02, 1.65815053e-01,
 1.53444767e-01, 1.97776090e-01, 2.33413022e-01, 6.95161339e-02,
 4.68937467e-02, 1.53353843e-01, 1.21391009e-01, 5.34892233e-02,
 7.17250734e-02, 5.68272472e-02, 1.87708648e-01, 8.92439810e-02,
 4.25453689e-02, 9.79721726e-02, 1.63975102e-02, 8.61687997e-02,
 6.31239429e-02, 3.34886298e-02, 1.09875888e-02, 3.83717041e-02,
 8.81504336e-02, 2.89832055e-01, 1.01621880e-01, 2.88982189e-01,
 5.81784635e-01, 3.29483923e-01, 8.00152896e-02, 7.31104074e-02,
 3.39216170e-02, 3.93543037e-02, 2.55189466e-02, 8.36160222e-02,
 5.48857747e-02, 8.78315482e-02, 1.10183077e-01, 1.37188728e-01,
 1.60562556e-01, 9.01524835e-02, 1.87554502e-01, 1.29971998e-01,
 3.71787543e-02, 6.31650566e-02, 1.65356465e-01, 1.58164919e-01,
 2.17681939e-01, 3.85248684e-01, 3.04221135e-01, 4.32125608e-01,
 2.94750935e-01, 1.22618052e-02, 2.74845469e-01, 2.10555482e-01,
 3.28604031e-03]),
array([0.          , 0.0078125, 0.015625 , 0.0234375, 0.03125  , 0.0390625,
 0.046875 , 0.0546875, 0.0625   , 0.0703125, 0.078125 , 0.0859375,
 0.09375  , 0.1015625, 0.109375 , 0.1171875, 0.125    , 0.1328125,
 0.140625 , 0.1484375, 0.15625  , 0.1640625, 0.171875 , 0.1796875,
 0.1875   , 0.1953125, 0.203125 , 0.2109375, 0.21875  , 0.2265625,
 0.234375 , 0.2421875, 0.25      , 0.2578125, 0.265625 , 0.2734375,
 0.28125  , 0.2890625, 0.296875 , 0.3046875, 0.3125   , 0.3203125,
 0.328125 , 0.3359375, 0.34375  , 0.3515625, 0.359375 , 0.3671875,
 0.375    , 0.3828125, 0.390625 , 0.3984375, 0.40625  , 0.4140625,
 0.421875 , 0.4296875, 0.4375   , 0.4453125, 0.453125 , 0.4609375,
 0.46875  , 0.4765625, 0.484375 , 0.4921875, 0.5       , 0.5078125,
```

```

0.515625 , 0.5234375, 0.53125 , 0.5390625, 0.546875 , 0.5546875,
0.5625 , 0.5703125, 0.578125 , 0.5859375, 0.59375 , 0.6015625,
0.609375 , 0.6171875, 0.625 , 0.6328125, 0.640625 , 0.6484375,
0.65625 , 0.6640625, 0.671875 , 0.6796875, 0.6875 , 0.6953125,
0.703125 , 0.7109375, 0.71875 , 0.7265625, 0.734375 , 0.7421875,
0.75 , 0.7578125, 0.765625 , 0.7734375, 0.78125 , 0.7890625,
0.796875 , 0.8046875, 0.8125 , 0.8203125, 0.828125 , 0.8359375,
0.84375 , 0.8515625, 0.859375 , 0.8671875, 0.875 , 0.8828125,
0.890625 , 0.8984375, 0.90625 , 0.9140625, 0.921875 , 0.9296875,
0.9375 , 0.9453125, 0.953125 , 0.9609375, 0.96875 , 0.9765625,
0.984375 , 0.9921875, 1. ]))

```



```

[27]: # transform the original dataset by differencing to attain stationarity and
      ↪ remove trend
diff_df = np.diff(df_teleco['Revenue'], axis=0)

```

```

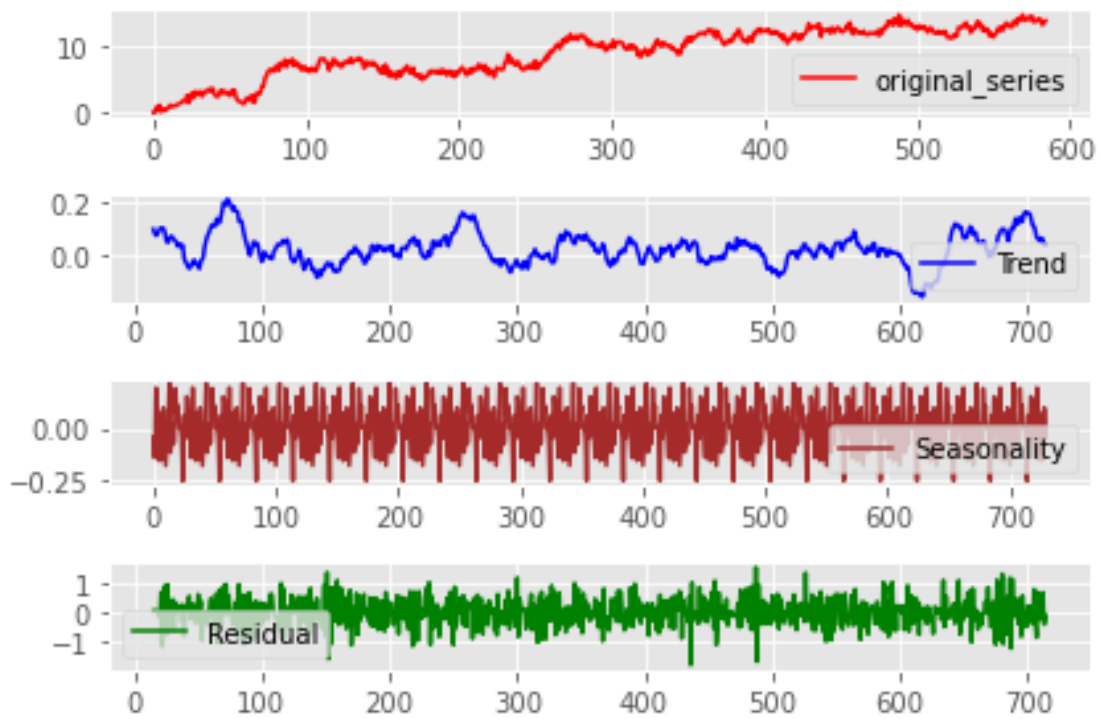
[28]: #decompose and plot again after differencing
decomposition=seasonal_decompose(diff_df, model='additive', period=30)
trend=decomposition.trend
seasonal=decomposition.seasonal
residual=decomposition.resid
plt.subplot(411)
plt.plot(df_train['Revenue'],color='red', label='original_series')
plt.legend(loc='best')

```

```

plt.subplot(412)
plt.plot(trend,color='blue', label='Trend')
plt.legend(loc='best')
plt.tight_layout()
plt.subplot(414)
plt.plot(residual,color='green', label='Residual')
plt.legend(loc='best')
plt.tight_layout()
plt.subplot(413)
plt.plot(seasonal,color='brown', label='Seasonality')
plt.legend(loc='best')
plt.tight_layout()
plt.show()


```




```

[29]: # Plot the trend component only again to confirm trend is eliminated
plt.figure(figsize=(12, 4))
plt.plot(trend, color='tab:blue')
plt.xlabel('Day')
plt.ylabel('Revenue Trend')
plt.title('Decomposed and differenced Revenue Trend')
plt.show()

```



```
[30]: #Plot the residuals( errors, basically zeros)
plt.plot(diff_df)
plt.show()
```



```
[31]: # Run adfuller test on the differenced DataFrame again to check for
      ↪ stationarity.
      #This time P-value is zero<0.05, hence data is stationary and ready for ARIMA

adft_diff = adfuller(diff_df, autolag='AIC')
print("1. ADF c-value : ",adft_diff[0])
```

```

print("2. P-Value : ", adft_diff[1])
print("3. Num Of Lags : ", adft_diff[2])
print("4. Num Of Observations Used :", adft_diff[3])
print("5. Critical Values :")
for key, val in adft[4].items():
    print("\t",key, ": ", val)

```

```


1. ADF c-value : -44.874527193875984
2. P-Value : 0.0
3. Num Of Lags : 0
4. Num Of Observations Used : 729
5. Critical Values :
    1% : -3.4393520240470554
    5% : -2.8655128165959236
    10% : -2.5688855736949163

```

```

[32]: #Test for stationarity again after differencing( confirmation, mean is_
      ↪virtually zero and constant)
diff_dff=pd.DataFrame(diff_df)
test_stationarity(diff_dff)

```



```

[33]: # Import auto_arima class
from pmdarima import auto_arima
from statsmodels.tsa.arima.model import ARIMA

```

[34]: *# running the auto ARIMA model to find the best model that minimizes AIC.*

```
arima_model = auto_arima(df_teleco['Revenue'], start_P=0,
                          start_q=0,
                          max_p=2,
                          max_q=2,
                          m=30,
                          seasonal=True,
                          d=1,
                          D=1,
                          trace=True,
                          error_action='ignore',
                          suppress_warnings=True,
                          stepwise=True)

arima_model.summary()
```

Performing stepwise search to minimize aic

ARIMA(2,1,0)(0,1,1)[30]	: AIC=inf, Time=212.11 sec
ARIMA(0,1,0)(0,1,0)[30]	: AIC=1617.639, Time=1.27 sec
ARIMA(1,1,0)(1,1,0)[30]	: AIC=1258.291, Time=10.41 sec
ARIMA(0,1,1)(0,1,1)[30]	: AIC=inf, Time=154.47 sec
ARIMA(1,1,0)(0,1,0)[30]	: AIC=1453.122, Time=2.28 sec
ARIMA(1,1,0)(2,1,0)[30]	: AIC=1151.386, Time=58.83 sec
ARIMA(1,1,0)(2,1,1)[30]	: AIC=inf, Time=281.73 sec
ARIMA(1,1,0)(1,1,1)[30]	: AIC=inf, Time=212.73 sec
ARIMA(0,1,0)(2,1,0)[30]	: AIC=1340.891, Time=31.82 sec
ARIMA(2,1,0)(2,1,0)[30]	: AIC=1153.381, Time=80.14 sec
ARIMA(1,1,1)(2,1,0)[30]	: AIC=1153.382, Time=87.75 sec
ARIMA(0,1,1)(2,1,0)[30]	: AIC=1194.259, Time=62.15 sec
ARIMA(2,1,1)(2,1,0)[30]	: AIC=1152.911, Time=169.81 sec
ARIMA(1,1,0)(2,1,0)[30] intercept	: AIC=1153.380, Time=278.64 sec

Best model: ARIMA(1,1,0)(2,1,0)[30]

Total fit time: 1647.332 seconds

[34]: <class 'statsmodels.iolib.summary.Summary'>

"""

SARIMAX Results

=====

Dep. Variable:	y	No. Observations:
731		
Model:	SARIMAX(1, 1, 0)x(2, 1, 0, 30)	Log Likelihood
-571.693		
Date:	Sat, 02 Apr 2022	AIC
1151.386		
Time:	01:31:21	BIC
1169.590		
Sample:	0	HQIC

1158.423

- 731

Covariance Type:

opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.4898	0.034	-14.216	0.000	-0.557	-0.422
ar.S.L30	-0.7098	0.037	-19.319	0.000	-0.782	-0.638
ar.S.L60	-0.4028	0.035	-11.425	0.000	-0.472	-0.334
sigma2	0.2915	0.015	19.261	0.000	0.262	0.321

===

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB):

0.89

Prob(Q): 0.99 Prob(JB):

0.64

Heteroskedasticity (H): 0.99 Skew:

0.05

Prob(H) (two-sided): 0.96 Kurtosis:

3.15

=====

===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

"""

```
[35]: # Build SARIMAX model on the train Data set using the (p,d,q)(P,D,Q)m results
      ↪from the model above
model_SAR = sm.tsa.SARIMAX(df_train['Revenue'], order=(1, 1, 0),
      ↪seasonal_order=(2, 1, 0, 30))
SARIMAX_Results = model_SAR.fit()

# Print results tables
print(SARIMAX_Results.summary())
```

SARIMAX Results

=====

=====

Dep. Variable:	Revenue	No. Observations:
585		
Model:	SARIMAX(1, 1, 0)x(2, 1, 0, 30)	Log Likelihood
-445.222		
Date:	Sat, 02 Apr 2022	AIC
898.443		
Time:	01:32:22	BIC
915.712		


```

Sample:                                0    HQIC
905.189

                                - 585
Covariance Type:                    opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          -0.4975      0.037     -13.290      0.000      -0.571      -0.424
ar.S.L30       -0.7178      0.042     -17.200      0.000      -0.800      -0.636
ar.S.L60       -0.4153      0.039     -10.662      0.000      -0.492      -0.339
sigma2         0.2815      0.016      17.165      0.000       0.249       0.314
=====
===
Ljung-Box (L1) (Q):                    0.06    Jarque-Bera (JB):
0.92
Prob(Q):                               0.80    Prob(JB):
0.63
Heteroskedasticity (H):                 0.98    Skew:
-0.03
Prob(H) (two-sided):                   0.87    Kurtosis:
3.19
=====
===

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

[36]: # call out forecast function
result_SAR = SARIMAX_Results.get_forecast()

```


```

[43]: # prediction on test set( 20% test, 80% train)

predictions = SARIMAX_Results.predict(585, 730, typ = 'levels').
↳rename('Predictions')

#Predict with respect to test set
plt.figure(figsize=(14, 6))
plt.plot(df_train['Revenue'], 'o-', color='black', label = '80% Train Dataset')
plt.plot(df_test['Revenue'], 'o-', color='blue', label = '20% Test Dataset')
plt.plot(predictions, 'o-', color='red', label = 'Predictions')
plt.title('Teleco Revenue Predictions')
plt.xlabel('Day')
plt.ylabel('Revenue, USD')
plt.legend(loc='best', fontsize = 8)
plt.show()

```



```
[44]: #Summary computations on test set
test_first= df_test['Revenue'].values.astype('float32')
forecast_test = result_SAR.predicted_mean
```

```
[78]: print('Expected : %.2f' % forecast_test)
print('Forecasted : %.2f' % test_first[0])
print('Standard Error : %.2f' % result_SAR.se_mean)
```

Expected : 13.68

Forecasted : 13.15

Standard Error : 0.53

```
[67]: # confident intervals
intervals = [0.2, 0.1, 0.05, 0.01]
for a in intervals:
    ci = result_SAR.conf_int(alpha=a)
    print('%0.1f%% Confidence Level: %.2f between %.2f and %.2f' %_
    →((1 - a) * 100, forecast_test, ci['lower Revenue'], ci['upper Revenue']))
ci
```

80.0% Confidence Level: 13.68 between 13.00 and 14.36

90.0% Confidence Level: 13.68 between 12.81 and 14.55

95.0% Confidence Level: 13.68 between 12.64 and 14.72

99.0% Confidence Level: 13.68 between 12.31 and 15.04

```
[67]:      lower Revenue  upper Revenue
585      12.311256      15.044417
```

```
[54]: # Run Mean Squared Error
MSE = mean_squared_error(df_test['Revenue'], predictions)
```

```

print('Summary')
print('MSE: ', round(MSE, 4))

# Run Root Mean Squared Error
RMSE = rmse(df_test['Revenue'], predictions)
print('RMSE: ', round(RMSE, 4))

```

Summary
MSE: 4.5606
RMSE: 2.1356

```

[40]: # make predictions with respect to the complete dataset
model00 = sm.tsa.statespace.SARIMAX(df_teleco['Revenue'], order=(1, 1, 0),
    ↪seasonal_order=(2, 1, 0, 30))
results00 = model00.fit()
# Print results tables
print(results00.summary())

```

SARIMAX Results

```

=====
=====
Dep. Variable:                Revenue    No. Observations:
731
Model:                SARIMAX(1, 1, 0)x(2, 1, 0, 30)    Log Likelihood
-571.693
Date:                Sat, 02 Apr 2022    AIC
1151.386
Time:                01:33:25    BIC
1169.590
Sample:                0    HQIC
1158.423
                                - 731
Covariance Type:                opg
=====
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.4898	0.034	-14.216	0.000	-0.557	-0.422
ar.S.L30	-0.7098	0.037	-19.319	0.000	-0.782	-0.638
ar.S.L60	-0.4028	0.035	-11.425	0.000	-0.472	-0.334
sigma2	0.2915	0.015	19.261	0.000	0.262	0.321

```

=====
=====
Ljung-Box (L1) (Q):                0.00    Jarque-Bera (JB):
0.89
Prob(Q):                0.99    Prob(JB):
0.64
Heteroskedasticity (H):                0.99    Skew:
0.05

```

Prob(H) (two-sided): 0.96 Kurtosis:
3.15


=====
===

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
[49]: # Forecast for the following quarter on the entire dataset( 2 years=731,
      ↪ 822=731+ 91, 91=quarter of 3 months)
forecast00 = results00.predict(731, 822, typ = 'level').rename('Teleco_
      ↪Forecast')

# plot predicted values
plt.figure(figsize=(14,6))
plt.plot(df_teleco['Revenue'], 'o-', color='black', label='Revenue Past')
plt.plot(forecast00, 'o-', color='red', label='Revenue Forecast' )
plt.legend(loc='best')
plt.show()
```



```
[50]: # Plot all for comparison
plt.figure(figsize=(14, 6))
plt.plot(df_train['Revenue'], 'o-', color='black', label = '80% Train Dataset')
plt.plot(df_test['Revenue'], 'o-', color='blue', label = '20% Test Dataset')
plt.plot(predictions, 'o-', color='green', label = 'Predictions on test set')
plt.plot(forecast00, 'o-', color='red', label='Revenue Forecast' )
plt.title('Teleco Revenue Predictions/forecast')
plt.xlabel('Day')
```

```
plt.ylabel('Revenue, USD')
plt.legend(loc='best', fontsize = 8)
plt.show()
```



```
[42]: #Try ARIMA model instead of Seasonal ARIMA
model_arima = ARIMA(df_teleco['Revenue'], order=(1,1,0))
results_ARIMA = model_arima.fit()
plt.plot(df_teleco['Revenue'], color='blue')
plt.plot(results_ARIMA.fittedvalues, color='red')
plt.show()


# summary of fit model
print('summary of fit model', results_ARIMA.summary())

# line plot of residuals
residuals = pd.DataFrame(results_ARIMA.resid)
residuals.plot()
plt.title('line plot of residuals')
plt.show()

# density plot of residuals
residuals.plot(kind='kde')
plt.title('density plot of residuals')
plt.show()

# summary stats of residuals
```

```
print( 'summary stats of residuals', residuals.describe())
```



summary of fit model

SARIMAX Results

```
=====
Dep. Variable:          Revenue    No. Observations:          731
Model:                ARIMA(1, 1, 0)  Log Likelihood            -490.355
Date:                 Sat, 02 Apr 2022  AIC                        984.710
Time:                 01:33:28    BIC                        993.896
Sample:               0          HQIC                        988.254
                   - 731
Covariance Type:      opg
=====
```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.4667	0.033	-14.213	0.000	-0.531	-0.402
sigma2	0.2243	0.013	17.782	0.000	0.200	0.249

===


```
Ljung-Box (L1) (Q):          0.00    Jarque-Bera (JB):
2.07
Prob(Q):                     0.98    Prob(JB):
0.36
Heteroskedasticity (H):      1.02    Skew:
-0.02
Prob(H) (two-sided):         0.89    Kurtosis:
2.74
```

```
=====
===
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).





```
summary stats of residuals      0
count  731.000000
mean    0.033076
std     0.472444
min     -1.593322
25%     -0.303294
50%      0.025629
75%      0.344243
max      1.257543
```

[]: